

Count-based differential expression analysis of RNA sequencing data using R and Bioconductor

Simon Anders¹, Davis J McCarthy^{2,3}, Yunshun Chen^{4,5}, Michal Okoniewski⁶, Gordon K Smyth^{4,7}, Wolfgang Huber¹ & Mark D Robinson^{8,9}

¹Genome Biology Unit, European Molecular Biology Laboratory, Heidelberg, Germany. ²Department of Statistics, University of Oxford, Oxford, UK. ³Wellcome Trust Centre for Human Genetics, University of Oxford, Oxford, UK. ⁴Bioinformatics Division, Walter and Eliza Hall Institute, Parkville, Victoria, Australia. ⁵Department of Medical Biology, University of Melbourne, Melbourne, Victoria, Australia. ⁶Functional Genomics Center UNI ETH, Zurich, Switzerland. ⁷Department of Mathematics and Statistics, University of Melbourne, Melbourne, Victoria, Australia. ⁸Institute of Molecular Life Sciences, University of Zurich, Zurich, Switzerland. ⁹SIB Swiss Institute of Bioinformatics, University of Zurich, Zurich, Switzerland. Correspondence should be addressed to M.D.R. (mark.robinson@imls.uzh.ch) or W.H. (whuber@embl.de).

Published online 22 August 2013; doi:10.1038/nprot.2013.099

RNA sequencing (RNA-seq) has been rapidly adopted for the profiling of transcriptomes in many areas of biology, including studies into gene regulation, development and disease. Of particular interest is the discovery of differentially expressed genes across different conditions (e.g., tissues, perturbations) while optionally adjusting for other systematic factors that affect the data-collection process. There are a number of subtle yet crucial aspects of these analyses, such as read counting, appropriate treatment of biological variability, quality control checks and appropriate setup of statistical modeling. Several variations have been presented in the literature, and there is a need for guidance on current best practices. This protocol presents a state-of-the-art computational and statistical RNA-seq differential expression analysis workflow largely based on the free open-source R language and Bioconductor software and, in particular, on two widely used tools, DESeq and edgeR. Hands-on time for typical small experiments (e.g., 4–10 samples) can be <1 h, with computation time <1 d using a standard desktop PC.

INTRODUCTION

Applications of the protocol

The RNA-seq platform^{1,2} addresses a multitude of applications, including relative expression analyses, alternative splicing, discovery of novel transcripts and isoforms, RNA editing, allele-specific expression and the exploration of non-model-organism transcriptomes.

Typically, tens of millions of sequences ('reads') are generated, and these, across several samples, form the starting point of this protocol. An initial and fundamental analysis goal is to identify genes whose expression level changes between conditions. In the simplest case, the aim is to compare expression levels between two conditions, e.g., stimulated versus unstimulated or wild type versus mutant. More complicated experimental designs can include additional experimental factors, potentially with multiple levels (e.g., multiple mutants, doses of a drug or time points) or may need to account for additional covariates (e.g. experimental batch or sex) or the pairing of samples (e.g., paired tumor and normal tissues from individuals). A crucial component of such an analysis is the statistical procedure used to call differentially expressed genes. This protocol covers two widely used tools for this task: DESeq³ and edgeR^{4–7}, both of which are available as packages of the Bioconductor software development project⁸.

Applications of these methods to biology and biomedicine are many. The methods described here are general and can be applied to situations in which observations are counts (typically, hundreds to tens of thousands of features of interest) and the goal is to discover changes in abundance. RNA-seq data are the standard use case (e.g., refs. 9,10), but many other differential analyses of counts are supported^{11,12}. For RNA-seq data, the strategy taken is to count the number of reads that fall into annotated genes and to perform statistical analysis on the table of counts to discover quantitative changes in expression levels between experimental groups. This counting approach is direct, flexible

and can be used for many types of count data beyond RNA-seq, such as comparative analysis of immunoprecipitated DNA^{11–14} (e.g., ChIP-seq, MBD-seq^{11,12}), proteomic spectral counts¹⁵ and metagenomics data.

Development of the protocol

Figure 1 gives the overall sequence of steps, from read sequences to feature counting to the discovery of differentially expressed genes, with a concerted emphasis on quality checks throughout. After initial checks on sequence quality, reads are mapped to a reference genome with a splice-aware aligner¹⁶; up to this point, this protocol^{3,6} is identical to many other pipelines (e.g., TopHat and Cufflinks¹⁷). From the set of mapped reads and either an annotation catalog or an assembled transcriptome, features, typically genes or transcripts, are counted and assembled into a table (rows for features and columns for samples). The statistical methods, which are integral to the differential expression discovery task, operate on a feature count table. Before the statistical modeling, further quality checks are encouraged to ensure that the biological question can be addressed. For example, a plot of sample relations can reveal possible batch effects and can be used to understand the similarity of replicates and the overall relationships between samples. After the statistical analysis of differential expression is carried out, a set of genes deemed to be differentially expressed or the corresponding statistics can be used in downstream interpretive analyses to confirm or generate further hypotheses.

Replication levels in designed experiments tend to be modest, often not much more than two or three. As a result, there is a need for statistical methods that perform well in small-sample situations. The low levels of replication rule out, for all practical purposes, distribution-free rank or permutation-based methods. Thus, for small-to-moderate sample sizes, the strategy used is to

Figure 1 | Count-based differential expression pipeline for RNA-seq data using edgeR and/or DESeq. Many steps are common to both tools, whereas the specific commands are different (Step 14). Steps within the edgeR or DESeq differential analysis can follow two paths, depending on whether the experimental design is simple or complex. Alternative entry points to the protocol are shown in orange boxes.

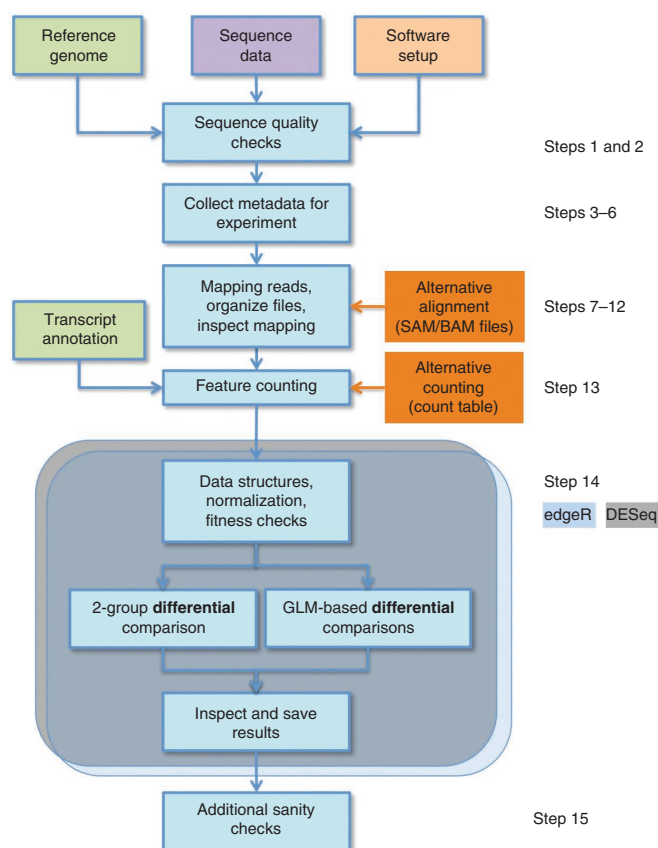
make formal distributional assumptions about the data observed. The advantage of parametric assumptions is the ability, through the wealth of existing statistical methodology, to make inferences about parameters of interest (i.e., changes in expression). For genome-scale count data, including RNA-seq, a convenient and well-established approximation is the negative binomial (NB) model (**Box 1**), which represents a natural extension of the Poisson model (i.e., a mixture of gamma-distributed rates) that was used in early studies¹⁸; notably, Poisson variation can only describe technical (i.e., sampling) variation.

To analyze differential expression, this protocol focuses on DESeq and edgeR, which implement general differential analyses on the basis of the NB model. These tools differ in their look and feel, and they estimate the dispersions differently but offer overlapping functionality (**Box 2**).

Variations and extensions of the protocol

This protocol presents a workflow built from a particular set of tools, but it is modular and extensible; thus, alternatives that offer special features (e.g., counting by allele) or additional flexibility (e.g., specialized mapping strategy) can be inserted as necessary. **Figure 1** highlights straightforward alternative entry points to the protocol (orange boxes). The count-based pipeline discussed here can be used in concert with other tools. For example, for species without an available well-annotated genome reference, Trinity¹⁹ or other assembly tools can be used to build a reference transcriptome; reads can then be aligned and counted, followed by the standard pipeline for differential analysis²⁰. Similarly, to perform differential analysis on novel genes in otherwise annotated genomes, the protocol could be expanded to include merged per-sample assemblies (e.g., Cuffmerge within the Cufflinks package^{17,21,22}) and used as input to counting tools.

The focus of this protocol is gene-level differential expression analysis. However, biologists are often interested in analyses



beyond that scope, and many possibilities now exist as extensions of the count-based framework discussed here. The full details of such analyses are not covered here, and we make only a sketch of some promising approaches. First, an obvious extension to gene-level counting is exon-level counting, given a catalog of transcripts. Reads can be assigned to the exons that they align to and be counted. Reads spanning exon-exon junctions can be counted at the junction level. The DEXSeq package uses a generalized linear model (GLM) that tests whether particular exons in a gene are preferentially used in a condition, over and above changes in gene-level expression. In edgeR, a similar strategy is taken, except that testing is done at the gene level by effectively asking whether

Box 1 | The NB model

The NB model has been shown to be a good fit to RNA-seq data⁷, yet it is flexible enough to account for biological variability. It provides a powerful framework (e.g., via GLMs) for analyzing arbitrarily complex experimental designs. NB models, as applied to genomic count data, make the assumption that an observation, say Y_{gj} (observed number of reads for gene g and sample j), has a mean μ_{gj} and a variance $\mu_{gj} + \phi_g \mu_{gj}^2$, where the dispersion $\phi_g > 0$ represents overdispersion relative to the Poisson distribution⁴. The mean parameters μ_{gj} depend on the sequencing depth for sample j as well as on the amount of RNA from gene g in the sample. Statistical procedures can be formulated to test for changes in expression level between experimental conditions, possibly adjusting for batch effects or other covariates, and to estimate the log-fold changes in expression.

The dispersion ϕ_g represents the squared coefficient of variation of the true expression levels between biologically independent RNA samples under the same experimental conditions, and hence the square root of ϕ_g is called the biological coefficient of variation⁷.

Obtaining good estimates of each gene's dispersion is critical for reliable statistical testing. Methods of estimating the genewise dispersion have received considerable attention^{3,4,31,59}. Unless the number of samples is large, stable estimation of the dispersion requires some sort of sharing of information between genes. One can average the variability across all genes⁵, or fit a global trend to the dispersion³, or can seek a more general compromise between individual gene and global dispersion estimators⁴.

Box 2 | Differences between DESeq and edgeR

The two packages described in this protocol, DESeq and edgeR, have similar strategies to perform differential analysis for count data. However, they differ in a few important areas. First, their look and feel differs. For users of the widely used limma package⁶⁰ (for analysis of microarray data), the data structures and steps in edgeR follow analogously. The packages differ in their default normalization: edgeR uses the trimmed mean of M values⁵⁶, whereas DESeq uses a relative log expression approach by creating a virtual library that every sample is compared against; in practice, the normalization factors are often similar. Perhaps most crucially, the tools differ in the choices made to estimate the dispersion. edgeR moderates feature-level dispersion estimates toward a trended mean according to the dispersion-mean relationship. In contrast, DESeq takes the maximum of the individual dispersion estimates and the dispersion-mean trend. In practice, this means DESeq is less powerful, whereas edgeR is more sensitive to outliers. Recent comparison studies have highlighted that no single method dominates another across all settings^{27,61,62}.

the exons are used proportionally across experiment conditions in the context of biological variation.

Comparison with other methods

Many tools exist for differential expression of counts, with slight variations of the method demonstrated in this protocol; these include, among others, baySeq²³, BBSseq²⁴, NOISeq²⁵ and QuasiSeq²⁶. The advantages and disadvantages of each tool are difficult to elicit for a given data set, but simulation studies show that edgeR and DESeq, despite the influx of many new tools, remain among the top performers²⁷.

The count-based RNA-seq analyses presented here consider the total output of a locus, without regard to the isoform diversity that may be present. This is of course a simplification. In certain situations, gene-level count-based methods may not recover true differential expression when some isoforms of a gene are upregulated and others are downregulated^{17,28}. Extensions of the gene-level count-based framework to differential exon usage are now available (e.g., DEXSeq²⁹). Recently, approaches have been proposed to estimate transcript-level expression and to build the uncertainty of these estimates into a differential analysis at the transcript level (e.g., BitSeq³⁰). Isoform deconvolution coupled with differential expression (e.g., Cuffdiff^{17,21,22}) is a plausible and popular alternative, but in general, isoform-specific expression estimation remains a difficult problem, especially if sequence reads are short, if genes whose isoforms overlap substantially are analyzed or if very deeply sequenced data are unavailable. At present, isoform deconvolution methods and transcript-level differential expression methods only support two-group comparisons. In contrast, counting is straightforward regardless of the configuration, and depth of data and arbitrarily complex experiments are naturally supported through GLMs (see **Box 3** for further details on feature counting). Recently, a flexible Bayesian framework for the analysis of 'random' effects in the context of GLM models and RNA-seq count data was made available in the ShrinkSeq package³¹. In addition, count-based methods that operate at the exon level, which share the NB framework, and flexible coverage-based methods have become available to address the limitations of gene-level analyses^{29,32,33}. These methods give a direct readout of differential exons, genes whose exons are used unequally or nonparallel coverage profiles, all of which reflect a change in isoform use.

Scope of this protocol

The aim of this protocol is to provide a concise workflow for a standard analysis, in a complete and easily accessible format, for users new to the field or to R. We describe a specific but very

common analysis task: the analysis of an RNA-seq experiment, comparing two groups of samples that differ in terms of their experimental treatment. We also cover one common complication: the need to account for a blocking factor.

In practice, users will need to adapt this pipeline to account for the circumstances of their experiment. More complicated experimental designs will require further considerations not covered here. Therefore, we emphasize that this protocol is not meant to replace the existing user guides, vignettes and online documentation for the packages and functions described. These provide a large body of information that is helpful for tackling tasks that go beyond the single-standard workflow presented here.

In particular, edgeR and DESeq have extensive user guides, downloadable from <http://www.bioconductor.org>, which cover a wide range of relevant topics. Please consult these comprehensive resources for further details. Another rich resource for answers to commonly asked questions is the Bioconductor mailing list (<http://bioconductor.org/help/mailling-list/>) as well as online resources such as seqanswers.com (<http://seqanswers.com/>), stackoverflow.com (<http://stackoverflow.com/>) and biostars.org (<http://www.biostars.org/>).

Multiple entry points to the protocol

As mentioned, this protocol is modular, in that users can use an alternative aligner or a different strategy (or software package) to count features. Two notable entry points (see orange boxes in **Fig. 1**) for the protocol include starting with either (i) a set of sequence alignment map (SAM)/binary alignment map (BAM) files from an alternative alignment algorithm or (ii) a table of counts. With SAM/BAM files in hand, users can start at Step 13, although it is often invaluable to carry along metadata information (Steps 3–6), postprocessing the alignment files may still be necessary (Step 9) and spot checks on the mapping are often useful (Steps 10–12). With a count table in hand, users can start at Step 14, where again the metadata information (Steps 3–6) will be needed for the statistical analysis. For users who wish to learn the protocol using the data analyzed here, the **Supplementary Data** gives an archive containing: the intermediate COUNT files used, a collated count table (counts) in CSV (comma-separated values) format, the metadata table (samples) in CSV format and the CSV file that was downloaded from the National Center for Biotechnology Information (NCBI)'s short read archive (SRA).

Experimental design

Replication. Some of the early RNA-seq studies were performed without biological replication. If the purpose of the experiment

Box 3 | Feature counting

In principle, counting reads that map to a catalog of features is straightforward. However, a few subtle decisions need to be made. For example, how should reads that fall within intronic regions (i.e., between two known exons) or beyond the annotated regions be counted? Ultimately, the answer to this question is guided by the chosen catalog that is presented to the counting software; depending on the protocol used, users should be conscious to include all features that are of interest, such as polyadenylated RNAs, small RNAs, long intergenic noncoding RNAs and so on. For simplicity and to avoid problems with mismatching chromosome identifiers and inconsistent coordinate systems, we recommend using the curated FASTA files and GTF files from Ensembl or the prebuilt indices packaged with GTF files from <http://tophat.cbcb.umd.edu/igenomes.shtml> whenever possible.

Statistical inference based on the NB distribution requires raw read counts as input. This is required to correctly model the Poisson component of the sample-to-sample variation. Therefore, it is crucial that units of evidence for expression are counted. No prior normalization or other transformation should be applied, including quantities such as RPKM (reads per kilobase model), FPKM (fragments per kilobase model) or otherwise depth-adjusted read counts. Both DESeq and edgeR internally keep the raw counts and normalization factors separate, as this full information is needed to correctly model the data. Notably, recent methods to normalize RNA-seq data for sample-specific G+C content effects use offsets that are presented to the GLM, while maintaining counts on their original scale^{63,64}.

Each paired-end read represents a single fragment of sequenced DNA, yet (at least) two entries for the fragment will appear in the corresponding BAM files. Some simplistic early methods that operated on BAM files considered these as separate entries, which led to overcounting and would ultimately overstate the significance of differential expression.

Typically, there will be reads that cannot be uniquely assigned to a gene, either because the read was aligned to multiple locations (multi-reads) or the read's position is annotated as part of several overlapping features. For the purpose of calling differential expression, such reads should be discarded. Otherwise, genuine differential expression of one gene might cause another gene to appear differentially expressed, erroneously, if reads from the first gene are counted for the second due to assignment ambiguity. In this protocol, we use the tool *htseq-count* of the Python package HTSeq, using the default union-counting mode; more details can be found at <http://www-huber.embl.de/users/anders/HTSeq/doc/count.html>. In addition, Bioconductor now offers various facilities for feature counting, including *easyRNASeq* in the *easyRNASeq* package⁶⁵, the *summarizeOverlaps* function in the *GenomicRanges* package and *qCount* in the *QuasR* (<http://www.bioconductor.org/packages/release/bioc/html/QuasR.html>) package.

is to make a general statement about a biological condition of interest (in statistical parlance, a population), for example, the effect of treating a certain cell line with a particular drug, then an experiment without replication is insufficient. Rapid developments in sequencing reduce technical variation but cannot possibly eliminate biological variability³⁴. Technical replicates are suited to studying properties of the RNA-seq platform¹⁶, but they do not provide information about the inherent biological variability in the system or the reproducibility of the biological result (for instance, its robustness to slight variations in cell density, passage number, drug concentration or medium composition). In other words, experiments without biological replication are suited to making a statement regarding one particular sample that existed on one particular day in one particular laboratory, but not whether anybody could reproduce this result. When no replicates are available, experienced analysts may still proceed, using one of the following options: (i) by performing a descriptive analysis with no formal hypothesis testing; (ii) by selecting a dispersion value on the basis of past experience; or (iii) by using housekeeping genes to estimate variability across all samples in the experiment.

In this context, it is helpful to remember the distinction between designed experiments, in which a well-characterized system (e.g., a cell line or a laboratory mouse strain) undergoes a fully controlled experimental procedure with minimal unintended variation, and observational studies, in which samples are often those of convenience (e.g., patients arriving at a clinic) that have been subjected to many uncontrolled environmental and genetic factors. Replication levels of two or three are often a practical

compromise between cost and benefit for designed experiments, but for observational studies, typically much larger group sizes (dozens or hundreds) are needed to reliably detect biologically meaningful results.

Confounding factors. In many cases, data are collected over time. In this situation, researchers should be mindful of factors that may unintentionally confound their results (e.g., batch effects), such as changes in reagent chemistry or software versions used to process their data³⁵. Users should make a concerted effort to reduce confounding effects through experimental design (e.g., randomization, blocking³⁶) and to keep track of versions, conditions (e.g., operators) of every sample, in the hope that these factors (or surrogates of them) can be differentiated from the biological factor(s) of interest in the downstream statistical modeling. In addition, there are emerging tools available that can discover and help eliminate unwanted variation in larger data sets^{37,38}, although these are relatively untested for RNA-seq data at present.

Software implementation. There are advantages to using a small number of software platforms for such a workflow, and these include simplified maintenance, training and portability. In principle, it is possible to do all computational steps in R and Bioconductor; however, for a few of the steps, the most mature and widely used tools are outside Bioconductor. Here R and Bioconductor are adopted to tie together the workflow and provide data structures, and their unique strengths in workflow components are leveraged, including statistical algorithms, visualization

Box 4 | Software versions

The original of this document was produced with Sweave⁶⁶ using the following versions of R and its packages:

```
> sessionInfo()
```

R output:

```
R version 3.0.0 (2013-04-03)
Platform: x86_64-unknown-linux-gnu (64-bit)

locale:
 [1] LC_CTYPE=en_CA.UTF-8      LC_NUMERIC=C               LC_TIME=en_CA.UTF-8
 [4] LC_COLLATE=en_CA.UTF-8    LC_MONETARY=en_CA.UTF-8    LC_MESSAGES=en_CA.UTF-8
 [7] LC_PAPER=C                LC_NAME=C                  LC_ADDRESS=C
[10] LC_TELEPHONE=C            LC_MEASUREMENT=en_CA.UTF-8 LC_IDENTIFICATION=C

attached base packages:
[1] parallel stats graphics grDevices utils datasets methods base

other attached packages:
[1] DESeq_1.12.0      locfit_1.5-9.1      Biobase_2.20.0      edgeR_3.2.3
[5] limma_3.16.2      ShortRead_1.18.0    latticeExtra_0.6-24 RColorBrewer_1.0-5
[9] Rsamtools_1.12.3  lattice_0.20-15     Biostrings_2.28.0    GenomicRanges_1.12.4
[13] IRanges_1.18.1    BiocGenerics_0.6.0  CacheSweave_0.6-1    stashR_0.3-5
[17] filehash_2.2-1

loaded via a namespace (and not attached):
[1] annotate_1.38.0  AnnotationDbi_1.22.5 bitops_1.0-5        DBI_0.2-7
[5] digest_0.6.3    genefilter_1.42.0    geneplotter_1.38.0 grid_3.0.0
[9] hwriter_1.3      RSQLite_0.11.3       splines_3.0.0       stats4_3.0.0
[13] survival_2.37-4 tools_3.0.0          XML_3.96-1.1        xtable_1.7-1
[17] zlibbioc_1.6.0
```

The versions of software packages used can be captured with the following commands (output is shown below each command):

```
> system("bowtie2 --version | grep align", intern=TRUE)
[1] "/usr/local/software/bowtie2-2.1.0/bowtie2-align version 2.1.0"

> system("tophat --version", intern=TRUE)
[1] "TopHat v2.0.8"

> system("htseq-count | grep version", intern=TRUE)
[1] "General Public License v3. Part of the 'HTSeq' framework, version 0.5.3p9."

> system("samtools 2 >&1 | grep Version", intern=TRUE)
[1] "Version: 0.1.18 (r982:295)"
```

and computation with annotation databases. Another major advantage of an R-based system, in terms of achieving best practices in genomic data analysis, is the opportunity for an interactive analysis whereby spot checks are made throughout the pipeline to guide the analyst. In addition, a wealth of tools is available for exploring, visualizing and cross-referencing genomic data. Although they are not used here directly, additional features of Bioconductor are readily available that will often be important for scientific projects that involve an RNA-seq analysis, including access to many different file formats, range-based computations, annotation resources, manipulation of sequence data and visualization.

In what follows, all Unix commands run at the command line appear in Courier font, prefaced by a dollar sign (\$):

```
$ my_unix_command
```

whereas R functions in the text appear as *myFunction*, and (typed) R input commands and output commands appear in **bold** and plain Courier font, respectively:

```
> x = 1:10
> median(x)
[1] 5.5
```

Note that in R, the operators = and <- can both be used for variable assignment (i.e., $z = 5$ and $z <- 5$ produce the same result, a new variable z with a numeric value). In this protocol, we use the = notation; in other places, users may also see the <- notation.

Constructing a metadata table (Steps 3–6). In general, we recommend starting from a sample metadata table that contains

sample identifiers, experimental conditions, blocking factors and file names. In our example, we construct this table from a file downloaded from the SRA. Users will often obtain a similar table from a local laboratory information management system (LIMS) or sequencing facility and can adapt this strategy to their own data sets.

Mapping reads to reference genome (Steps 7 and 8). In the protocol, R is used to tie the pipeline together (i.e., loop through the set of samples and construct the full tophat2 command), with the hope of reducing typing and copy-and-paste errors. Many alternatives and variations are possible: users can use R to create and call the tophat2 commands, to create the commands (and call tophat2 independently of a Unix shell), or to assemble the commands manually independently of R. tophat2 creates a directory for each sample with the mapped reads in a BAM file, called `accepted_hits.bam`. Note that BAM files, and equivalently SAM files (an uncompressed text version of BAM), are the de facto standard file for alignments. Therefore, alternative mapping tools that produce BAM/SAM files could be inserted into the protocol at this point.

Organizing BAM and SAM files (Step 9). The set of files containing mapped reads (from tophat2, `accepted_hits.bam`) (typically) needs to be transformed before it can be used with other downstream tools. In particular, the samtools command is used to prepare variations of the mapped reads. Specifically, a sorted and indexed version of the BAM file is created, which can be used in genome browsers such as IGV; a sorted-by-name SAM file is created, which is compatible with the feature-counting software of htseq-count. Alternative feature-counting tools (e.g., in Bioconductor) may require different inputs.

Design matrix. For more complex designs (i.e., beyond two-group comparisons), users need to provide a design matrix that specifies the factors that are expected to affect expression levels. As mentioned above, GLMs can be used to analyze arbitrarily complex experiments, and the design matrix is the means by which the experimental design is described mathematically, including both biological factors of interest and other factors not of direct interest, such as batch effects. For example, Section 4.5 of the edgeR User's Guide ('RNA-seq of pathogen inoculated *Arabidopsis* with batch effects') or Section 4 of the DESeq vignette ('Multi-factor designs') presents worked case studies with batch effects. The design matrix is central for such complex differential expression analyses, and users may wish to consult a linear modeling textbook³⁹ or a local statistician to make sure their design matrix is appropriately specified.

Reproducible research. We recommend that users keep a record of all commands (R and Unix) and the software versions used in their analysis so that other researchers (e.g., collaborators, reviewers) can reproduce the results (**Box 4**). In practice, this is best achieved by keeping the complete transcript of the computer commands interweaved with the textual narrative in a single, executable document⁴⁰. R provides many tools to facilitate the authoring of executable documents, including the *Sweave* function and the knitr package. The *sessionInfo* function helps with documenting package versions and related information. A recent integration with Rstudio is rpubs.com (<http://rpubs.com/>), which provides seamless integration of 'mark-down' text with R commands for easy web-based display. For language-independent authoring, a powerful tool is provided by Emacs org-mode.

MATERIALS

EQUIPMENT

▲ **CRITICAL** For many of the software packages listed below, new features and optimizations are constantly developed and released, so we highly recommend using the most recent stable version as well as reading the (corresponding) documentation for the version used. The package versions used in the production of this article are given in **Box 4**

Operating system

• This protocol assumes users have a Unix-like operating system (i.e., Linux or MacOS X), with a bash shell or similar. All commands given here are meant to be run in a terminal window. Although it is possible to follow this protocol with a Microsoft Windows machine (e.g., using the Unix-like Cygwin; <http://www.cygwin.com/>), the additional steps required are not discussed here

Software

- An aligner to map short reads to a genome that is able to deal with reads that straddle introns¹⁶. The aligner tophat2 (refs. 21,41) is illustrated here, but others, such as GSNAP⁴², SpliceMap⁴³, Subread⁴⁴ or STAR⁴⁵, can be used
- (Optional) A tool to visualize alignment files, such as the Integrated Genome Viewer (IGV⁴⁶, or Savant^{47,48}). IGV is a Java tool with 'web start' (downloadable from <http://www.broadinstitute.org/software/igv/log-in>), i.e., it can be started from a web browser and needs no

explicit installation at the operating system level, provided a Java Runtime Environment is available

- The R statistical computing environment, downloadable from <http://www.r-project.org/>
- A number of Bioconductor⁸ packages, specifically ShortRead⁴⁹, DESeq³ and edgeR^{6,7}, and possibly GenomicRanges, GenomicFeatures and org.Dm.eg.db, as well as their dependencies
- The samtools program⁵⁰ (<http://samtools.sourceforge.net/>) for manipulation of SAM- and BAM-formatted files
- The HTSeq package (<http://www-huber.embl.de/users/anders/HTSeq/doc/overview.html>) for counting of mapped reads
- (Optional) If users wish to work with data from the SRA, they will need the SRA Toolkit (<http://www.ncbi.nlm.nih.gov/Traces/sra/sra.cgi?cmd=show&f=software&m=software&s=software>)

Input file formats

- In general, the starting point is a collection of FASTQ files, the commonly used format for reads from Illumina sequencing machines. The modifications necessary for mapping reads from other platforms are not discussed here

Example data

- The data set published by Brooks *et al.*⁵¹ is used here to demonstrate the workflow. This data set consists of seven RNA-seq samples, each a cell

culture of *Drosophila melanogaster* S2 cells. Three samples were treated with siRNA targeting the splicing factor *pasilla* (CG1844) ('knockdown') and four samples are untreated ('control'). Our aim is to identify genes that change in expression between knockdown and control. Brooks *et al.*⁵¹ have sequenced some of their libraries in single-end mode and others in paired-end mode. This allows us to demonstrate two variants of the workflow: if we ignore the differences in library type, the samples only differ by their experimental condition (knockdown or control), and the analysis is a simple comparison between two sample groups. We refer to this setting as an experiment with a simple design. If we want to account for library type as a blocking factor, our samples differ in more than one aspect (i.e., we have a complex design). To deal with the latter scenario, we use edgeR and DESeq's functions to fit GLMs.

EQUIPMENT SETUP

Install bowtie2, tophat2 and samtools Download and install samtools from <http://samtools.sourceforge.net>. bowtie2 and tophat2 have binary versions available for Linux and Mac OS X platforms. These can be downloaded from <http://bowtie-bio.sourceforge.net/index.shtml> and <http://tophat.cbcb.umd.edu/>. Consult the documentation on those sites for further information if necessary.

Install R and required Bioconductor packages Download the latest release version of R from <http://cran.r-project.org/> and install it. Consult the R Installation and Administration manual if necessary. A useful quick reference for R commands can be found at <http://cran.r-project.org/doc/contrib/Short-refcard.pdf>. To install Bioconductor packages, start R by issuing the command R in a terminal window and type:

```
> source("http://www.Bioconductor.org/biocLite.R")
> biocLite("BiocUpgrade")
> biocLite(c("ShortRead", "DESeq", "edgeR"))
```

This retrieves an automatic installation tool (*biocLite*) and installs the version-matched packages. In addition, the installation tool will automatically download and install all other prerequisite packages. Versions of Bioconductor packages are matched to versions of R. Hence, to use current versions of Bioconductor packages, it is necessary to use a current version of R. Note that R and Bioconductor maintain a stable release version and a

development version at all times. Unless a special need exists for a particular new functionality, users should use the release version.

Download the example data To download SRA repository data, an automated process may be desirable. For example, from <http://www.ncbi.nlm.nih.gov/sra?term=SRP001537> (the entire experiment corresponding to GEO accession GSE18508), users can download a table of the metadata into a comma-separated tabular file 'SraRunInfo.csv' (see the **Supplementary Data**, which contains an archive of various files used in this protocol). To do this, click on 'Send to:' (top right corner), select 'File', select format 'RunInfo' and click on 'Create File'. Read this CSV file 'SraRunInfo.csv' into R, and select the subset of samples that we are interested in (using R's string matching function `grep`), corresponding to the 22 SRA files shown in Figure 2 by:

```
> sri = read.csv("SraRunInfo.csv",
                stringsAsFactors=FALSE)
> keep = grep("CG8144|Untreated-",
              sri$ LibraryName)
> sri = sri[keep,]
```

The following R commands automate the download of the 22 SRA files to the current working directory (the functions *getwd* and *setwd* can be used to retrieve and set the working directory, respectively):

```
> fs = basename(sri$download_path)
> for(i in 1:nrow(sri))
    download.file(sri$download_path[i], fs[i])
```

▲ CRITICAL This download is only required if data originate from the SRA. Brooks *et al.*⁵¹ deposited their data in the SRA of the NCBI's Gene Expression Omnibus (GEO)⁵² under accession number GSE18508 and a subset of this data set is used here to illustrate the pipeline. Specifically, SRA files corresponding to the 4 'Untreated' (control) and 3 'CG8144 RNAi' (knockdown) samples need to be downloaded.

Alternative download tools The R-based download of files described above is just one way to capture several files in a semiautomatic fashion. Users can alternatively use the batch tools *wget* (Unix/Linux) or *curl* (Mac OS X), or

Figure 2 | Screenshot of Metadata available from SRA.

download using a web browser. The (truncated) verbose output of the above R download commands looks as follows:

```
trying URL 'ftp://ftp-private.ncbi.nlm.nih.gov/sra/sra-instant/reads/ByRun/sra/SRR/SRR031/SRR031714/SRR031714.sra'
ftp data connection made, file length 415554366 bytes
opened URL
=====
downloaded 396.3 Mb
trying URL 'ftp://ftp-private.ncbi.nlm.nih.gov/sra/sra-instant/reads/ByRun/sra/SRR/SRR031/SRR031715/SRR031715.sra'
ftp data connection made, file length 409390212 bytes
opened URL
=====
downloaded 390.4 Mb
[... truncated ...]
```

Convert SRA to FASTQ format Typically, sequencing data from a sequencing facility will come in (compressed) FASTQ format. The SRA, however, uses its own, compressed, SRA format. In order to convert the example data to FASTQ, use the `fastq-dump` command from the SRA Toolkit on each SRA file. Note that the use of R's `system` command is just one possibility. Users may choose to type the 22 `fastq-dump` commands manually into the Unix shell rather than using R to construct them. R can be used to construct the required shell commands, starting from the 'SraRunInfo.csv' metadata table, as follows:

```
> stopifnot( all(file.exists(fs)) ) # assure FTP
  download was successful
> for(f in fs) {
  cmd = paste("fastq-dump --split-3", f)
  cat(cmd, "\n")
  system(cmd) # invoke command
}
```

Using the `cat` command It is not absolutely necessary to use `cat` to print out the current command, but it serves the purpose of knowing what is currently running in the shell:

```
fastq-dump --split-3 SRR031714.sra
Written 5327425 spots for SRR031714.sra
Written 5327425 spots total
fastq-dump --split-3 SRR031715.sra
Written 5248396 spots for SRR031715.sra
Written 5248396 spots total
[... truncated ...]
```

▲ CRITICAL Be sure to use the `--split-3` option, which splits mate-pair reads into separate files. After this command, single and paired-end data will produce one and two FASTQ files, respectively. For paired-end data, the file names will be suffixed `_1.FASTQ` and `_2.FASTQ`; otherwise, a single file with the extension `.FASTQ` will be produced.

Download the reference genome Download the reference genome sequence for the organism under study in (compressed) FASTA format. Some useful resources, among others, include: the general Ensembl FTP server (<http://www.ensembl.org/info/data/ftp/index.html>), the Ensembl plants FTP server (<http://plants.ensembl.org/info/data/ftp/index.html>), the Ensembl metazoa FTP server (<http://metazoa.ensembl.org/info/data/ftp/index.html>) and the University of California Santa Cruz (UCSC) current genomes FTP server (<ftp://hgdownload.cse.ucsc.edu/goldenPath/currentGenomes/>).

Using Ensembl For Ensembl, choose the 'FASTA (DNA)' link instead of 'FASTA (cDNA)', as alignments to the genome, not the transcriptome, are desired. For *D. melanogaster*, the file labeled 'toplevel' combines all chromosomes. Do not use the 'repeat-masked' files (indicated by 'rm' in the file name), as handling repeat regions should be left to the alignment algorithm. The *Drosophila* reference genome can be downloaded from Ensembl and uncompressed using the following Unix commands:

```
$ wget ftp://ftp.ensembl.org/pub/release-70/
  fasta/drosophila_melanogaster/dna/Drosophila_
  melanogaster.BDGP5.70.dna.toplevel.fa.gz
$ gunzip Drosophila_melanogaster.BDGP5.70.dna.
  toplevel.fa.gz
```

For genomes provided by UCSC, users can select their genome of interest, proceed to the 'bigZips' directory and download the 'chromFa.tar.gz'; as above, this could be done using the `wget` command. Note that `bowtie2` and `tophat2` indices for many commonly used reference genomes can be downloaded directly from <http://tophat.cbcb.umd.edu/igenomes.shtml>.

Get gene model annotations Download a gene transfer format (GTF) file with gene models for the organism of interest. For species covered by Ensembl, the Ensembl FTP site mentioned above contains links to such files. The gene model annotation for *D. melanogaster* can be downloaded and uncompressed using:

```
$ wget ftp://ftp.ensembl.org/pub/release-70/gtf/
  drosophila_melanogaster/Drosophila_melanogaster.
  BDGP5.70.gtf.gz
$ gunzip Drosophila_melanogaster.BDGP5.70.gtf.gz
```

▲ CRITICAL Make sure that the gene annotation uses the same coordinate system as the reference FASTA file. Here, both files use BDGP5 (i.e., release 5 of the assembly provided by the Berkeley *Drosophila* Genome Project), as is apparent from the file names. To be on the safe side, here, we recommend always downloading the FASTA reference sequence and the GTF annotation data from the same resource provider.

▲ CRITICAL As an alternative, the UCSC Table Browser (<http://genome.ucsc.edu/cgi-bin/hgTables>) can be used to generate GTF files on the basis of a selected annotation (e.g., RefSeq genes). However, at the time of writing, GTF files obtained from the UCSC Table Browser do not contain correct gene IDs, which can cause problems with downstream tools such as `htseq-count`, unless they are corrected manually.

Build the reference index Before reads can be aligned, the reference FASTA files need to be preprocessed into an index that allows the aligner easy access. To build a `bowtie2`-specific index from the FASTA file mentioned above, use the command:

```
$ bowtie2-build -f Drosophila_melanogaster.
  BDGP5.70.dna.toplevel.fa Dmel_BDGP5_70
```

A set of BT2 files will be produced, with names starting with `Dmel_BDGP5_70` specified above. This procedure needs to be run only once for each reference genome used. As mentioned, pre-built indices for many commonly used genomes are available from <http://tophat.cbcb.umd.edu/igenomes.shtml>.

PROCEDURE

Assess sequence quality control with ShortRead ● TIMING ~2 h

1| At the R prompt, type the commands (you may first need to use *setwd* to set the working directory to where the FASTQ files are situated):

```
> library("ShortRead")
> fqQC = qa(dirPath=".", pattern=".fastq$", type="fastq")
> report(fqQC, type="html", dest="fastQQAreport")
```

? TROUBLESHOOTING

2| Use a web browser to inspect the generated HTML file (here, stored in the 'fastQQAreport' directory) with the quality-assessment report (see ANTICIPATED RESULTS for further details)

Collect metadata of experimental design ● TIMING <1 h

3| Create a table of metadata called 'samples' (see 'Constructing metadata table' in Experimental Design). This step needs to be adapted for each data set, and many users may find a spreadsheet program useful for this step, from which data can be imported into the table *samples* by the *read.csv* function. For our example data, we chose to construct the *samples* table programmatically from the table of SRA files.

4| Collapse the initial table (*sri*) to one row per sample:

```
> sri$LibraryName = gsub("S2_DRSC_", "", sri$LibraryName) # trim label
> samples = unique(sri[,c("LibraryName", "LibraryLayout")])
> for(i in seq_len(nrow(samples))) {
  rw = (sri$LibraryName == samples$LibraryName[i])
  if(samples$LibraryLayout[i] == "PAIRED") {
    samples$fastq1[i] = paste0(sri$Run[rw], "_1.fastq", collapse=",")
    samples$fastq2[i] = paste0(sri$Run[rw], "_2.fastq", collapse=",")
  } else {
    samples$fastq1[i] = paste0(sri$Run[rw], ".fastq", collapse=",")
    samples$fastq2[i] = ""
  }
}
```

5| Add important or descriptive columns to the metadata table (experimental groupings are set on the basis of the 'LibraryName' column, and a label is created for plotting):

```
> samples$condition = "CTL"
> samples$condition[grep("RNAi", samples$LibraryName)] = "KD"
> samples$shortname = paste(substr(samples$condition, 1, 2),
                             substr(samples$LibraryLayout, 1, 2),
                             seq_len(nrow(samples)), sep=".")
```

6| As the downstream statistical analysis of differential expression relies on this table, carefully inspect (and correct, if necessary) the metadata table. In particular, verify that there exists one row per sample, that all columns of information are populated and that the file names, labels and experimental conditions are correct.

```
> samples
```

PROTOCOL

R output:

	LibraryName	Library Layout	fastq1	fastq2	condition	shortname
1	Untreated-3	PAIRED	SRR031714_1.fastq, ...	SRR031714_2.fastq, ...	CTL	CT.PA.1
2	Untreated-4	PAIRED	SRR031716_1.fastq, ...	SRR031716_2.fastq, ...	CTL	CT.PA.2
3	CG8144_RNAi-3	PAIRED	SRR031724_1.fastq, ...	SRR031724_2.fastq, ...	KD	KD.PA.3
4	CG8144_RNAi-4	PAIRED	SRR031726_1.fastq, ...	SRR031726_2.fastq, ...	KD	KD.PA.4
5	Untreated-1	SINGLE	SRR031708.fastq, ...		CTL	CT.SI.5
6	CG8144_RNAi-1	SINGLE	SRR031718.fastq, ...		KD	KD.SI.6
7	Untreated-6	SINGLE	SRR031728.fastq, ...		CTL	CT.SI.7

Align the reads (using tophat2) to the reference genome ● **TIMING** ~45 min per sample

7| By using R string manipulation, construct the Unix commands to call tophat2. Given the metadata table samples, it is convenient to use R to create the list of shell commands, as follows:

```
> gf = "Drosophila_melanogaster.BDGP5.70.gtf"
> bowind = "Dme1_BDGP5_70"
> cmd = with(samples, paste("tophat2 -G", gf, "-p 5 -o",
                             LibraryName, bowind, fastq1, fastq2))
> cmd
```

R output:

```
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o Untreated-3 Dme1_BDGP5_70 \
SRR031714_1.fastq,SRR031715_1.fastq SRR031714_2.fastq,SRR031715_2.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o Untreated-4 Dme1_BDGP5_70 \
SRR031716_1.fastq,SRR031717_1.fastq SRR031716_2.fastq,SRR031717_2.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o CG8144_RNAi-3 Dme1_BDGP5_70 \
SRR031724_1.fastq,SRR031725_1.fastq SRR031724_2.fastq,SRR031725_2.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o CG8144_RNAi-4 Dme1_BDGP5_70 \
SRR031726_1.fastq,SRR031727_1.fastq SRR031726_2.fastq,SRR031727_2.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o Untreated-1 Dme1_BDGP5_70 \
SRR031708.fastq,SRR031709.fastq,SRR031710.fastq,SRR031711.fastq,SRR031712.fastq,
SRR031713.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o CG8144_RNAi-1 Dme1_BDGP5_70 \
SRR031718.fastq,SRR031719.fastq,SRR031720.fastq,SRR031721.fastq,SRR031722.fastq,
SRR031723.fastq
tophat2 -G Drosophila_melanogaster.BDGP5.70.gtf -p 5 -o Untreated-6 Dme1_BDGP5_70 \
SRR031728.fastq,SRR031729.fastq
```

▲ **CRITICAL STEP** In the call to tophat2, the option -G points tophat2 to a GTF file of annotation to facilitate mapping reads across exon-exon junctions (some of which can be found de novo), -o specifies the output directory, -p specifies the number of threads to use (this may affect run times and can vary depending on the resources available). Other parameters can be specified here, as needed; see the appropriate documentation for the tool and version you are using. The first argument, Dme1_BDGP5_70 is the name of the index (built in advance), and the second argument is a list of all FASTQ files with reads for the sample. Note that the FASTQ files are concatenated with commas, without spaces. For experiments with paired-end reads, pairs of FASTQ files are given as separate arguments and the order in both arguments must match.

? TROUBLESHOOTING

8| Run these commands (i.e., copy and paste) in a Unix terminal.

▲ **CRITICAL STEP** Many similar possibilities exist for this step (see 'Experimental design' for further details). Users can use the R function *system* to execute these commands direct from R, cut-and-paste the commands into a separate Unix shell or

store the list of commands in a text file and use the Unix 'source' command. In addition, users could construct the Unix commands independently of R.

? TROUBLESHOOTING

Organize, sort and index the BAM files and create SAM files ● TIMING ~1 h

9| Organize the BAM files into a single directory, sort and index them and create SAM files by running the following R-generated commands:

```
> for(i in seq_len(nrow(samples))) {
  lib = samples$LibraryName[i]
  ob = file.path(lib, "accepted_hits.bam")

  # sort by name, convert to SAM for htseq-count
  cat(paste0("samtools sort -n ",ob," ",lib,"_sn"),"\n")
  cat(paste0("samtools view -o ",lib,"_sn.sam ",lib,"_sn.bam"),"\n")

  # sort by position and index for IGV
  cat(paste0("samtools sort ",ob," ",lib,"_s"),"\n")
  cat(paste0("samtools index ",lib,"_s.bam"),"\n\n")
}
```

R output:

```
samtools sort -n Untreated-3/accepted_hits.bam Untreated-3_sn
samtools view -o Untreated-3_sn.sam Untreated-3_sn.bam
samtools sort Untreated-3/accepted_hits.bam Untreated-3_s
samtools index Untreated-3_s.bam

samtools sort -n Untreated-4/accepted_hits.bam Untreated-4_sn
samtools view -o Untreated-4_sn.sam Untreated-4_sn.bam
samtools sort Untreated-4/accepted_hits.bam Untreated-4_s
samtools index Untreated-4_s.bam

samtools sort -n CG8144_RNAi-3/accepted_hits.bam CG8144_RNAi-3_sn
samtools view -o CG8144_RNAi-3_sn.sam CG8144_RNAi-3_sn.bam
samtools sort CG8144_RNAi-3/accepted_hits.bam CG8144_RNAi-3_s
samtools index CG8144_RNAi-3_s.bam

samtools sort -n CG8144_RNAi-4/accepted_hits.bam CG8144_RNAi-4_sn
samtools view -o CG8144_RNAi-4_sn.sam CG8144_RNAi-4_sn.bam
samtools sort CG8144_RNAi-4/accepted_hits.bam CG8144_RNAi-4_s
samtools index CG8144_RNAi-4_s.bam

samtools sort -n Untreated-1/accepted_hits.bam Untreated-1_sn
samtools view -o Untreated-1_sn.sam Untreated-1_sn.bam
samtools sort Untreated-1/accepted_hits.bam Untreated-1_s
samtools index Untreated-1_s.bam

samtools sort -n CG8144_RNAi-1/accepted_hits.bam CG8144_RNAi-1_sn
samtools view -o CG8144_RNAi-1_sn.sam CG8144_RNAi-1_sn.bam
samtools sort CG8144_RNAi-1/accepted_hits.bam CG8144_RNAi-1_s
samtools index CG8144_RNAi-1_s.bam
```





Figure 3 | Screenshot of reads aligning across exon junctions.

```
samtools sort -n Untreated-6/accepted_hits.bam Untreated-6_sn
samtools view -o Untreated-6_sn.sam Untreated-6_sn.bam
samtools sort Untreated-6/accepted_hits.bam Untreated-6_s
samtools index Untreated-6_s.bam
```

▲ CRITICAL STEP Users should be conscious of the disk space that may get used in these operations. In the command above, sorted-by-name SAM and BAM files (for htseq-count), as well as a sorted-by-chromosome-position BAM file (for IGV), are created for each original accepted_hits.bam file. User may wish to delete (some of) these intermediate files after the steps below.

Inspect alignments with IGV ● TIMING <20 min

10| Start IGV, select the *D. melanogaster* (dm3) genome, and then load the BAM files (with s in the filename) as well as the GTF file.

11| Zoom in on an expressed transcript until individual reads are shown and check whether the reads align at and across exon-exon junctions, as expected, given the annotation (**Fig. 3**).

12| If any positive and negative controls are known for the system under study (e.g., known differential expression), direct the IGV browser to these regions to confirm that the relative read density is different according to expectation.

Count reads using htseq-count ● TIMING ~3 h

13| Add the names of the COUNT files to the metadata table and call HTSeq from the following R-generated Unix commands:

```
> samples$countf = paste(samples$LibraryName, "count", sep=".")
> gf = "Drosophila_melanogaster.BDGP5.70.gtf"
> cmd = paste0("htseq-count -s no -a 10 ", samples$LibraryName,
  "_sn.sam ", gf, " > ", samples$countf)
> cmd
```

R output:

```
htseq-count -s no -a 10 Untreated-3_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > Untreated-3.count
```

```
htseq-count -s no -a 10 Untreated-4_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > Untreated-4.count

htseq-count -s no -a 10 CG8144_RNAi-3_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > CG8144_RNAi-3.count

htseq-count -s no -a 10 CG8144_RNAi-4_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > CG8144_RNAi-4.count

htseq-count -s no -a 10 Untreated-1_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > Untreated-1.count

htseq-count -s no -a 10 CG8144_RNAi-1_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > CG8144_RNAi-1.count

htseq-count -s no -a 10 Untreated-6_sn.sam \
Drosophila_melanogaster.BDGP5.70.gtf > Untreated-6.count
```

▲ **CRITICAL STEP** The option `-s` signifies that the data are not from a stranded protocol (this may vary by experiment) and the `-a` option specifies a minimum score for the alignment quality.

? TROUBLESHOOTING

14| For differential expression analysis with edgeR, follow option A for simple designs and option B for complex designs; for differential expression analysis with DESeq, follow option C for simple designs and option D for complex designs.

(A) edgeR—simple design

(i) Load the edgeR package and use the utility function, *readDGE*, to read in the COUNT files created from htseq-count:

```
> library("edgeR")
> counts = readDGE(samples$countf)$counts
```

? TROUBLESHOOTING

(ii) Filter weakly expressed and noninformative (e.g., non-aligned) features using a command like:

```
> noint = rownames(counts) %in%
      c("no_feature", "ambiguous", "too_low_aQual",
        "not_aligned", "alignment_not_unique")
> cpms = cpm(counts)
> keep = rowSums(cpms > 1) >= 3 & !noint
> counts = counts[keep,]
```

▲ **CRITICAL STEP** In edgeR, it is recommended to remove features without at least 1 read per million in *n* of the samples, where *n* is the size of the smallest group of replicates (here, *n* = 3 for the knockdown group).

(iii) Visualize and inspect the count table as follows:

```
> colnames(counts) = samples$shortname
> head( counts[,order(samples$condition)], 5 )
```

R output:

	CT.PA.1	CT.PA.2	CT.SI.5	CT.SI.7	KD.PA.3	KD.PA.4	KD.SI.6
FBgn0000008	76	71	137	82	87	68	115
FBgn0000017	3498	3087	7014	3926	3029	3264	4322
FBgn0000018	240	306	613	485	288	307	528
FBgn0000032	611	672	1479	1351	694	757	1361
FBgn0000042	40048	49144	97565	99372	70574	72850	95760

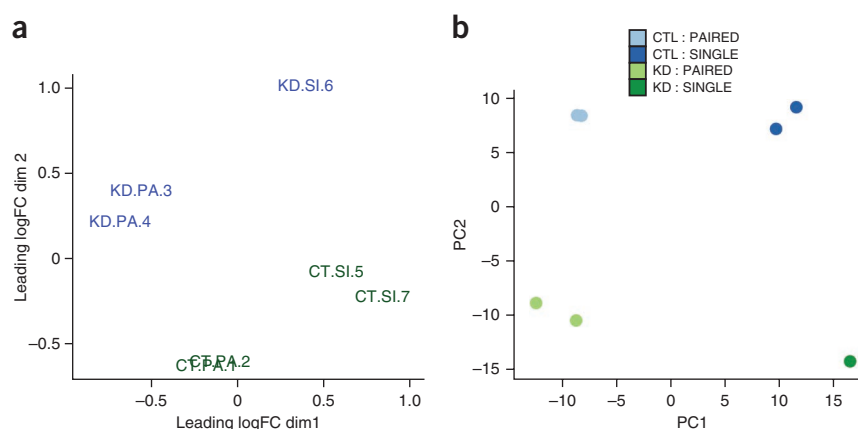


Figure 4 | Plots of sample relations. (a) By using a count-specific distance measure, edgeR's *plotMDS* produces a multidimensional scaling plot showing the relationship between all pairs of samples. (b) DESeq's *plotPCA* makes a principal component (PC) plot of VST (variance-stabilizing transformation)-transformed count data. CT or CTL, control; KD, knockdown.

(iv) Create a *DGEList* object (edgeR's container for RNA-seq count data), as follows:

```
> d = DGEList(counts=counts, group=samples$condition)
```

(v) Estimate normalization factors using:

```
> d = calcNormFactors(d)
```

(vi) Inspect the relationships between samples using a multidimensional scaling (MDS) plot, as shown in **Figure 4**:

```
> plotMDS(d, labels=samples$shortname,
           col=c("darkgreen", "blue")[factor(samples$condition)])
```

(vii) Estimate tagwise dispersion (simple design) using:

```
> d = estimateCommonDisp(d)
> d = estimateTagwiseDisp(d)
```

(viii) Create a visual representation of the mean-variance relationship using the *plotMeanVar* (**Fig. 5a**) and *plotBCV* (**Fig. 5b**) functions, as follows:

```
> plotMeanVar(d, show.tagwise.vars=TRUE, NBline=TRUE)
> plotBCV(d)
```

(ix) Test for differential expression ('classic' edgeR), as follows:

```
> de = exactTest(d, pair=c("CTL", "KD"))
```

(x) Follow Step 14B(vi–ix).

(B) edgeR—complex design

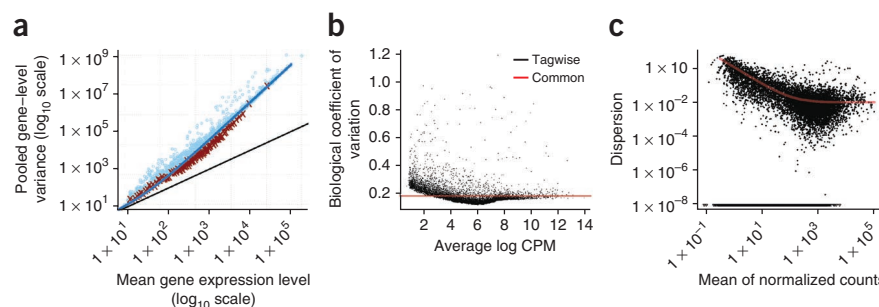
(i) Follow Step 14A(i–vi).

(ii) Create a design matrix (see 'Experimental design' for further details) to specify the factors that are expected to affect expression levels:

```
> design = model.matrix(~ LibraryLayout + condition, samples)
> design
```

R output:

Figure 5 | Plots of mean-variance relationship and dispersion. (a) edgeR's *plotMeanVar* can be used to explore the mean-variance relationship; each dot represents the estimated mean and variance for each gene, with binned variances as well as the trended common dispersion overlaid. (b) edgeR's *plotBCV* illustrates the relationship of biological coefficient of variation versus mean log CPM. (c) DESeq's *plotDispEsts* shows the fit of dispersion versus mean. CPM, counts per million.



```
(Intercept) LibraryLayoutSINGLE conditionKD
1 1 0 0
2 1 0 0
3 1 0 1
4 1 0 1
5 1 1 0
6 1 1 1
7 1 1 0
```

```
attr("assign")
[1] 0 1 2
attr("contrasts")
attr("contrasts")$LibraryLayout
[1] "contr.treatment"
attr("contrasts")$condition
[1] "contr.treatment"
```

(iii) Estimate dispersion values, relative to the design matrix, using the Cox-Reid (CR)-adjusted likelihood^{7,53}, as follows:

```
> d2 = estimateGLMTrendedDisp(d, design)
> d2 = estimateGLMTagwiseDisp(d2, design)
```

(iv) Given the design matrix and dispersion estimates, fit a GLM to each feature:

```
> f = glmFit(d2, design)
```

(v) Perform a likelihood ratio test, specifying the difference of interest (here, knockdown versus control, which corresponds to the third column of the above design matrix):

```
> de = glmLRT(f, coef=3)
```


(vi) Use the *topTags* function to present a tabular summary of the differential expression statistics (note that *topTags* operates on the output of *exactTest* or *glmLRT*, but only the latter is shown here):

```
> tt = topTags(de, n=nrow(d))
> head(tt$table)
```

R output:

	logFC	logCPM	LR	PValue	FDR
FBgn0039155	-4.61	5.87	902	3.96e-198	2.85e-194
FBgn0025111	2.87	6.86	641	2.17e-141	7.81e-138
FBgn0039827	-4.05	4.40	457	2.11e-101	5.07e-98
FBgn0035085	-2.58	5.59	408	9.31e-91	1.68e-87
FBgn0000071	2.65	4.73	365	2.46e-81	3.54e-78
FBgn0003360	-3.12	8.42	359	3.62e-80	4.34e-77

Figure 6 | M ('minus') versus A ('add') plots for RNA-seq data. (a) edgeR's *plotSmear* function plots the log-fold change (i.e., the log ratio of normalized expression levels between two experimental conditions) against the log counts per million (CPM). (b) Similarly, DESeq's *plotMA* displays differential expression (log-fold changes) versus expression strength (log average read count).



(vii) Inspect the depth-adjusted reads per million for some of the top differentially expressed genes:

```
> nc = cpm(d, normalized.lib.sizes=TRUE)
> rn = rownames(tt$table)
> head(nc[rn,order(samples$condition)],5)
```

R output:

	CT.PA.1	CT.PA.2	CT.SI.5	CT.SI.7	KD.PA.3	KD.PA.4	KD.SI.6
FBgn0039155	91.07	98.0	100.75	106.78	3.73	4.96	3.52
FBgn0025111	34.24	31.6	26.64	28.46	247.43	254.28	188.39
FBgn0039827	39.40	36.7	30.09	34.47	1.66	2.77	2.01
FBgn0035085	78.06	81.4	63.59	74.08	13.49	14.13	10.99
FBgn0000071	9.08	9.2	7.48	5.85	52.08	55.93	45.65

(viii) Create a graphical summary, such as an M (log-fold change) versus A (log-average expression) plot⁵⁴, here showing the genes selected as differentially expressed (with a 5% false discovery rate; **Fig. 6**):

```
> deg = rn[tt$table$FDR < .05]
> plotSmear(d, de.tags=deg)
```

(ix) Save the result table as a CSV file (alternative formats are possible) as follows:

```
> write.csv(tt$table, file="toptags_edgeR.csv")
```

(C) DESeq—simple design

(i) Create a *data.frame* with the required metadata, i.e., the names of the count files and experimental conditions. Here we derive it from the samples table created in Step 3.

```
> samplesDESeq = with(samples,
  data.frame(shortname = I(shortname), countf = I(countf),
    condition = condition,
    LibraryLayout = LibraryLayout))
```

(ii) Load the DESeq package and create a *CountDataSet* object (DESeq's container for RNA-seq data) from the count tables and corresponding metadata:

```
> library("DESeq")
> cds = newCountDataSetFromHTSeqCount(samplesDESeq)
```

? TROUBLESHOOTING

(iii) Estimate normalization factors using:

```
> cds = estimateSizeFactors(cds)
```

(iv) Inspect the size factors using:

```
> sizeFactors(cds)
```

R output:

```
CT.PA.1 CT.PA.2 KD.PA.3 KD.PA.4 CT.SI.5 KD.SI.6 CT.SI.7
0.699    0.811    0.822    0.894    1.643    1.372    1.104
```

- (v) To inspect sample relationships, invoke a variance-stabilizing transformation and inspect a principal component analysis (PCA) plot (**Fig. 4b**):

```
> cdsB = estimateDispersions(cds, method="blind")
> vsd = varianceStabilizingTransformation(cdsB)
> p = plotPCA(vsd, intgroup=c("condition", "LibraryLayout"))
```

- (vi) Use *estimateDispersions* to calculate dispersion values:

```
> cds = estimateDispersions(cds)
```

- (vii) Inspect the estimated dispersions using the *plotDispEsts* function (**Fig. 5c**), as follows:

```
> plotDispEsts(cds)
```

- (viii) Perform the test for differential expression by using *nbinomTest*, as follows:

```
> res = nbinomTest(cds, "CTL", "KD")
```

- (ix) Given the table of differential expression results, use *plotMA* to display differential expression (log-fold changes) versus expression strength (log-average read count), as follows (**Fig. 6b**):

```
> plotMA(res)
```

- (x) Inspect the result tables of significantly upregulated and downregulated genes, at a 10% false discovery rate (FDR) as follows:

```
> resSig = res[which(res$padj < 0.1),]
> head( resSig[ order(resSig$log2FoldChange, decreasing=TRUE), ] )
```

R output:

	id	baseMean	baseMeanA	baseMeanB	foldChange	log2FoldChange	pval	padj
1515	FBgn0013696	1.46	0.000	3.40	Inf	Inf	4.32e-03	6.86e-02
13260	FBgn0085822	1.93	0.152	4.29	28.2	4.82	4.54e-03	7.11e-02
13265	FBgn0085827	8.70	0.913	19.08	20.9	4.39	1.02e-09	9.57e-08
15470	FBgn0264344	3.59	0.531	7.68	14.5	3.86	4.55e-04	1.10e-02
8153	FBgn0037191	4.43	0.715	9.39	13.1	3.71	5.35e-05	1.78e-03
1507	FBgn0013688	23.82	4.230	49.95	11.8	3.56	3.91e-21	1.38e-18

```
> head( resSig[ order(resSig$log2FoldChange, decreasing=FALSE), ] )
```

R output:

	id	baseMean	baseMeanA	baseMeanB	foldChange	log2FoldChange	pval	padj
13045	FBgn0085359	60.0	102.2	3.78	0.0370	-4.76	7.65e-30	4.88e-27
9499	FBgn0039155	684.1	1161.5	47.59	0.0410	-4.61	3.05e-152	3.88e-148
2226	FBgn0024288	52.6	88.9	4.25	0.0478	-4.39	2.95e-32	2.09e-29
9967	FBgn0039827	246.2	412.0	25.08	0.0609	-4.04	1.95e-82	8.28e-79
6279	FBgn0034434	104.5	171.8	14.72	0.0856	-3.55	8.85e-42	9.40e-39
6494	FBgn0034736	203.9	334.9	29.38	0.0877	-3.51	6.00e-41	5.88e-38

- (xi) Count the number of genes with significant differential expression at a FDR of 10%:

```
> table( res$padj < 0.1 )
```

R output:

```
FALSE    TRUE
11861    885
```

- (xii) Create a persistent storage of results using, for example, a CSV file:

```
> write.csv(res, file="res_DESeq.csv")
```

- (xiii) Perform a sanity check by inspecting a histogram of unadjusted *P* values (Fig. 7) for the differential expression results, as follows:

```
> hist(res$pval, breaks=100)
```




Figure 7 | Histogram of *P* values from gene-by-gene statistical tests.

(D) DESeq—complex design

- (i) Follow Step 14C(i–v).

- (ii) Calculate the CR-adjusted profile likelihood⁵³ dispersion estimates relative to the factors specified, developed by McCarthy *et al.*⁷, according to:

```
> cds = estimateDispersions(cds, method = "pooled-CR",
                             modelFormula = count ~ LibraryLayout + condition)
```

- (iii) Test for differential expression in the GLM setting by fitting both a full model and reduced model (i.e., with the factor of interest taken out):

```
> fit1 = fitNbinomGLMs(cds, count ~ LibraryLayout + condition)
> fit0 = fitNbinomGLMs(cds, count ~ LibraryLayout)
```

- (iv) By using the two fitted models, compute likelihood ratio statistics and associated *P* values, as follows:

```
> pval = nbinomGLMTest(fit1, fit0)
```

- (v) Adjust the reported *P* values for multiple testing:

```
> padj = p.adjust(pval, method="BH")
```

- (vi) Assemble a result table from full model fit and the raw and adjusted *P* values and print the first few upregulated and downregulated genes (FDR <10%):

```
> res = cbind(fit1, pval=pval, padj=padj)
> resSig = res[which(res$padj < 0.1),]
> head( resSig[ order(resSig$conditionKD, decreasing=TRUE), ] )
```

R output:

```
(Intercept) LibraryLayoutSINGLE conditionKD deviance converged pval padj
FBgn0013696 -70.96 36.829 37.48 5.79e-10 TRUE 3.52e-03 5.51e-02
FBgn0085822 -5.95 4.382 5.14 2.07e+00 TRUE 4.72e-03 7.07e-02
FBgn0085827 -3.96 4.779 5.08 2.89e+00 TRUE 5.60e-03 8.05e-02
FBgn0264344 -2.59 2.506 4.26 6.13e-01 TRUE 3.86e-04 9.17e-03
FBgn0261673 3.53 0.133 3.37 1.39e+00 TRUE 0.00e+00 0.00e+00
FBgn0033065 2.85 -0.421 3.03 4.07e+00 TRUE 8.66e-15 1.53e-12

> head( resSig[ order(resSig$conditionKD, decreasing=FALSE), ] )
```

R output:

	(Intercept)	LibraryLayoutSINGLE	conditionKD	deviance	converged	pval	padj
FBgn0031923	1.01	1.2985	-32.26	1.30	TRUE	0.00528	0.077
FBgn0085359	6.37	0.5782	-4.62	3.32	TRUE	0.00000	0.000
FBgn0039155	10.16	0.0348	-4.62	3.39	TRUE	0.00000	0.000
FBgn0024288	6.71	-0.4840	-4.55	1.98	TRUE	0.00000	0.000
FBgn0039827	8.79	-0.2272	-4.06	2.87	TRUE	0.00000	0.000
FBgn0034736	8.54	-0.3123	-3.57	2.09	TRUE	0.00000	0.000

(vii) Follow Step 14C(xi–xiii).

15 | As another spot check, point the IGV genome browser (with GTF and BAM files loaded) to a handful of the top differentially expressed genes and confirm that the counting and differential expression statistics are appropriately represented.

? TROUBLESHOOTING

Troubleshooting advice can be found in **Table 1**.

TABLE 1 | Troubleshooting table.

Step	Problem	Possible reason	Solution
1, 14A(i), 14C(ii)	An error occurs when loading a Bioconductor package	Version mismatch	Make sure the most recent version of R is installed; reinstall packages using <i>biocLite</i>
7, 8	An error occurs while mapping reads to reference genome	Wrong files made available or version mismatch	Carefully check the command submitted, the documentation for the aligner and the setup steps (e.g., building an index); check that there is no clash between bowtie and bowtie2
13	An error occurs counting features	GTF format violation	Use an Ensembl GTF format or coerce your file into a compatible format. In particular, verify that each line of type exon contains attributes named gene_id and transcript_id, and ensure that their values are correct
14	Errors in fitting statistical models or running statistical tests	Wrong inputs, outdated version of software	Ensure versions of R and Bioconductor packages are up to date and check the command issued; if a command is correct but an error persists, post a message to the Bioconductor mailing list (http://bioconductor.org/help/ mailing-list/) according to the posting guide (http://bioconductor.org/help/ mailing-list/posting-guide/)

● TIMING

Running this protocol on the SRA-downloaded data will take <10 h on a machine with eight cores and 8 GB of RAM; with a machine with more cores, mapping of different samples can be run simultaneously. The time is largely spent on quality checks of reads, read alignment and feature counting; computation time for the differential expression analysis is comparatively smaller.

Step 1, sequence quality checks: ~2 h

Steps 3–6, organizing metadata: <1 h

Steps 7 and 8, read alignment: ~45 min per sample

Step 9, organize, sort and index the BAM files and create SAM files: ~1 h

Steps 10–12, inspect alignments with IGV: <20 min

Step 13, feature counting: ~3 h

Step 14, differential analysis: variable; computational time is often <20 min

Step 15, additional spot checks: <20 min

ANTICIPATED RESULTS

Sequencing quality checks

Step 1 results in an HTML report for all included FASTQ files. Users should inspect these (Step 2) and look for persistence of low-quality scores, over-representation of adapter sequence and other potential problems. From these inspections, users may choose to remove low-quality samples, trim ends of reads (e.g., using FASTX; http://hannonlab.cshl.edu/fastx_toolkit/) or modify alignment parameters. Note that a popular non-Bioconductor alternative for sequencing quality checks is FastQC (<http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>).

Feature counting

In Step 13, we used htseq-count for feature counting. The output is a COUNT file (two columns: identifier, count) for each sample. Many alternatives exist inside and outside of Bioconductor to arrive at a table of counts given BAM (or SAM) files and a set of features (e.g., from a GTF file); see **Box 3** for further considerations. Each cell in the count table will be an integer that indicates how many reads in the sample overlap with the respective feature. Non-informative rows, such as features that are not of interest or those that have low overall counts, can be filtered. We recommend removing rows with a low overall sum of counts (or low CPM), as this generally increases the statistical power of the differential expression analysis⁵⁵.

Normalization

As different libraries will be sequenced to different depths, offsets are built in the statistical model to ensure that parameters are comparable. The term normalization is often used for that, but it should be noted that the raw read counts are not actually altered⁵⁶. By default, edgeR uses the number of mapped reads (i.e., count table column sums) and estimates an additional normalization factor to account for sample-specific effects (e.g., diversity)⁵⁶; these two factors are combined and used as an offset in the NB model. Analogously, DESeq defines a virtual reference sample by taking the median of each gene's values across samples and then computes size factors as the median of ratios of each sample to the reference sample. Generally, the ratios of the size factors should roughly match the ratios of the library sizes. Dividing each column of the count table by the corresponding size factor yields normalized count values, which can be scaled to give a counts per million interpretation (see also edgeR's *cpm* function). From an M (log ratio) versus A (log expression strength) plot, count data sets typically show a (left-facing) trombone shape, reflecting the higher variability of log ratios at lower counts (**Fig. 6**). In addition, points will typically be centered around a log ratio of 0 if the normalization factors are calculated appropriately, although this is just a general guide.

Sample relations

The quality of the sequencing reactions (Step 1) themselves is only a part of the quality assessment procedure. In Steps 14A(vi) or 14C(v), a 'fitness for use'⁵⁷ check is performed (relative to the biological question of interest) on the count data before statistical modeling. edgeR adopts a straightforward approach that compares the relationship between all pairs of samples, using a count-specific pairwise distance measure (i.e., biological coefficient of variation) and an MDS plot for visualization (**Fig. 4a**). Analogously, DESeq performs a variance-stabilizing transformation and explores sample relationships using a PCA plot (**Fig. 4b**). In either case, the analysis for the current data set highlights that library type (single-end or paired-end) has a systematic effect on the read counts and provides an example of a data-driven modeling decision: here, a GLM-based analysis that accounts for the (assumed linear) effect of library type jointly with the biological factor of interest (i.e., knockdown versus control) is recommended. In general, users should be conscious that the degree of variability between the biological replicates (e.g., in an MDS or PCA plot) will ultimately effect the calling of differential expression. For example, a single outlying sample may drive increased dispersion estimates and compromise the discovery of differentially expressed features. No general prescription is available for when and whether to delete outlying samples.

Dispersion estimation

As mentioned above, getting good estimates of the dispersion parameter is critical to the inference of differential expression. For simple designs, edgeR uses the quantile-adjusted conditional maximum (weighted) likelihood estimator^{4,5}, whereas DESeq uses a method-of-moments estimator³. For complex designs, the dispersion estimates are made relative to the design matrix, using the CR-adjusted likelihood^{7,53}; both DESeq and edgeR use this estimator. edgeR's estimates are always moderated toward a common trend, whereas DESeq chooses the maximum of the individual estimate and a smooth fit (dispersion versus mean) over all genes. A wide range of dispersion-mean relationships exist in RNA-seq data, as viewed by edgeR's *plotBCV* or DESeq's *plotDispEsts*; case studies with further details are presented in both edgeR's and DESeq's user guides.

Differential expression analysis

DESeq and edgeR differ slightly in the format of results outputted, but each contains columns for log-fold change (log), counts per million (or mean by condition), likelihood ratio statistic (for GLM-based analyses), as well as raw and adjusted

P values. By default, *P* values are adjusted for multiple testing using the Benjamini-Hochberg⁵⁸ procedure. If users enter tabular information to accompany the set of features (e.g., annotation information), edgeR has a facility to carry feature-level information into the results table.

Post-differential analysis sanity checks

Figure 7 (Step 14C(xiii)) shows the typical features of a *P* value histogram resulting from a good data set: a sharp peak at the left side, containing genes with strong differential expression, a ‘floor’ of values that are approximately uniform in the interval [0, 1], corresponding to genes that are not differentially expressed (for which the null hypothesis is true), and a peak at the upper end at 1, resulting from discreteness of the NB test for genes with overall low counts. The latter component is often less pronounced, or even absent, when the likelihood ratio test is used. In addition, users should spot check genes called as differentially expressed by loading the sorted BAM files into a genome browser.

Note: Supplementary information is available in the [online version of the paper](#).

ACKNOWLEDGMENTS We thank X. Zhou for comparing counting methods, O. Nikolayeva for feedback on an earlier version of the manuscript and participants in the European Conference on Computational Biology Workshop (Basel, September 2012) for their feedback. G.K.S. acknowledges funding from a National Health and Medical Research Council Project Grant (no. 1023454). D.J.M. acknowledges funding from the General Sir John Monash Foundation, Australia. M.D.R. wishes to acknowledge funding from the University of Zurich's Research Priority Program in Systems Biology and Functional Genomics and Swiss National Science Foundation Project grant (no. 143883). S.A., W.H. and M.D.R. acknowledge funding from the European Commission through the 7th Framework Collaborative Project RADIANT (grant agreement no. 305626).

AUTHOR CONTRIBUTIONS S.A. and W.H. are authors of the DESeq package. D.J.M., Y.C., G.K.S. and M.D.R. are authors of the edgeR package. S.A., M.O., W.H. and M.D.R. initiated the protocol format on the basis of the ECCB 2012 Workshop. S.A. and M.D.R. wrote the first draft and additions were made from all authors.

COMPETING FINANCIAL INTERESTS The authors declare no competing financial interests.

Reprints and permissions information is available online at <http://www.nature.com/reprints/index.html>.

- Mortazavi, A., Williams, B.A., McCue, K., Schaeffer, L. & Wold, B. Mapping and quantifying mammalian transcriptomes by RNA-seq. *Nat. Methods* **5**, 621–628 (2008).
- Wang, Z., Gerstein, M. & Snyder, M. RNA-seq: a revolutionary tool for transcriptomics. *Nat. Rev. Genet.* **10**, 57–63 (2009).
- Anders, S. & Huber, W. Differential expression analysis for sequence count data. *Genome Biol.* **11**, R106 (2010).
- Robinson, M.D. & Smyth, G.K. Moderated statistical tests for assessing differences in tag abundance. *Bioinformatics* **23**, 2881–2887 (2007).
- Robinson, M.D. & Smyth, G.K. Small-sample estimation of negative binomial dispersion, with applications to SAGE data. *Biostatistics* **9**, 321–332 (2008).
- Robinson, M.D., McCarthy, D.J. & Smyth, G.K. edgeR: a Bioconductor package for differential expression analysis of digital gene expression data. *Bioinformatics* **26**, 139–140 (2010).
- McCarthy, D.J., Chen, Y. & Smyth, G.K. Differential expression analysis of multifactor RNA-seq experiments with respect to biological variation. *Nucleic Acids Res.* **40**, 4288–4297 (2012).
- Gentleman, R.C. *et al.* Bioconductor: open software development for computational biology and bioinformatics. *Genome Biol.* **5**, R80 (2004).
- Zemach, A. *et al.* The *Arabidopsis* nucleosome remodeler DDM1 allows DNA methyltransferases to access H1-containing heterochromatin. *Cell* **153**, 193–205 (2013).
- Lam, M.T. *et al.* Rev-Erbs repress macrophage gene expression by inhibiting enhancer-directed transcription. *Nature* **498**, 511–515 (2013).
- Ross-Innes, C.S. *et al.* Differential oestrogen receptor binding is associated with clinical outcome in breast cancer. *Nature* **481**, 389–393 (2012).
- Robinson, M.D. *et al.* Copy-number-aware differential analysis of quantitative DNA sequencing data. *Genome Res.* **22**, 2489–2496 (2012).
- Vanharanta, S. *et al.* Epigenetic expansion of VHL-HIF signal output drives multiorgan metastasis in renal cancer. *Nat. Med.* **19**, 50–56 (2013).
- Samstein, R.M. *et al.* Foxp3 exploits a pre-existent enhancer landscape for regulatory T cell lineage specification. *Cell* **151**, 153–166 (2012).
- Johnson, E.K. *et al.* Proteomic analysis reveals new cardiac-specific dystrophin-associated proteins. *PLoS ONE* **7**, e43515 (2012).
- Fonseca, N.A., Rung, J., Brazma, A. & Marioni, J.C. Tools for mapping high-throughput sequencing data. *Bioinformatics* **28**, 3169–3177 (2012).
- Trapnell, C. *et al.* Differential gene and transcript expression analysis of RNA-seq experiments with TopHat and Cufflinks. *Nat. Protoc.* **7**, 562–578 (2012).
- Bullard, J.H., Purdom, E., Hansen, K.D. & Dudoit, S. Evaluation of statistical methods for normalization and differential expression in mRNA-seq experiments. *BMC Bioinform.* **11**, 94 (2010).
- Grabherr, M.G. *et al.* Full-length transcriptome assembly from RNA-seq data without a reference genome. *Nat. Biotechnol.* **29**, 644–652 (2011).
- Siebert, S. *et al.* Differential gene expression in the siphonophore *Nanomia bijuga* (Cnidaria) assessed with multiple next-generation sequencing workflows. *PLoS ONE* **6**, 12 (2011).
- Trapnell, C., Pachter, L. & Salzberg, S.L. TopHat: discovering splice junctions with RNA-seq. *Bioinformatics* **25**, 1105–1111 (2009).
- Trapnell, C. *et al.* Transcript assembly and quantification by RNA-seq reveals unannotated transcripts and isoform switching during cell differentiation. *Nat. Biotechnol.* **28**, 511–515 (2010).
- Hardcastle, T.J. & Kelly, K.A. baySeq: empirical Bayesian methods for identifying differential expression in sequence count data. *BMC Bioinform.* **11**, 422 (2010).
- Zhou, Y.-H., Xia, K. & Wright, F.A. A powerful and flexible approach to the analysis of RNA sequence count data. *Bioinformatics* **27**, 2672–2678 (2011).
- Tarazona, S., Garcia-Alcalde, F., Dopazo, J., Ferrer, A. & Conesa, A. Differential expression in RNA-seq: a matter of depth. *Genome Res.* **21**, 2213–2223 (2011).
- Lund, S.P., Nettleton, D., McCarthy, D.J. & Smyth, G.K. Detecting differential expression in RNA-sequence data using quasi-likelihood with shrunken dispersion estimates. *Stat. Appl. Genet. Mol. Biol.* **11**, pii (2012).
- Soneson, C. & Delorenzi, M. A comparison of methods for differential expression analysis of RNA-seq data. *BMC Bioinform.* **14**, 91 (2013).
- Lareau, L.F., Inada, M., Green, R.E., Wengrod, J.C. & Brenner, S.E. Unproductive splicing of SR genes associated with highly conserved and ultraconserved DNA elements. *Nature* **446**, 926–929 (2007).
- Anders, S., Reyes, A. & Huber, W. Detecting differential usage of exons from RNA-seq data. *Genome Res.* **22**, 2008–2017 (2012).
- Glaus, P., Honkela, A. & Rattray, M. Identifying differentially expressed transcripts from RNA-seq data with biological variation. *Bioinformatics* **28**, 1721–1728 (2012).
- Van De Wiel, M.A. *et al.* Bayesian analysis of RNA sequencing data by estimating multiple shrinkage priors. *Biostatistics* **14**, 113–128 (2013).

32. Blekhman, R., Marioni, J.C., Zumbo, P., Stephens, M. & Gilad, Y. Sex-specific and lineage-specific alternative splicing in primates. *Genome Res.* **20**, 180–189 (2010).
33. Okoniewski, M.J. *et al.* Preferred analysis methods for single genomic regions in RNA sequencing revealed by processing the shape of coverage. *Nucleic Acids Res.* **40**, e63 (2012).
34. Hansen, K.D., Wu, Z., Irizarry, R.A. & Leek, J.T. Sequencing technology does not eliminate biological variability. *Nat. Biotechnol.* **29**, 572–573 (2011).
35. Leek, J.T. *et al.* Tackling the widespread and critical impact of batch effects in high-throughput data. *Nat. Rev. Genet.* **11**, 733–739 (2010).
36. Auer, P.L. & Doerge, R.W. Statistical design and analysis of RNA sequencing data. *Genetics* **185**, 405–416 (2010).
37. Gagnon-Bartsch, J.A. & Speed, T.P. Using control genes to correct for unwanted variation in microarray data. *Biostatistics* **13**, 539–552 (2011).
38. Leek, J.T. & Storey, J.D. Capturing heterogeneity in gene expression studies by surrogate variable analysis. *PLoS Genet.* **3**, 1724–1735 (2007).
39. Myers, R.M. *Classical and Modern Regression with Applications* 2nd edn. (Duxbury Classic Series, 2000).
40. Gentleman, R. Reproducible research: a bioinformatics case study. *Stat. Appl. Genet. Mol. Biol.* **4**, Article2 (2005).
41. Trapnell, C. & Salzberg, S.L. How to map billions of short reads onto genomes. *Nat. Biotechnol.* **27**, 455–457 (2009).
42. Wu, T.D. & Nacu, S. Fast and SNP-tolerant detection of complex variants and splicing in short reads. *Bioinformatics* **26**, 873–881 (2010).
43. Wang, K. *et al.* MapSplice: Accurate mapping of RNA-seq reads for splice junction discovery. *Nucleic Acids Res.* **38**, e178 (2010).
44. Liao, Y., Smyth, G.K. & Shi, W. The Subread aligner: fast, accurate and scalable read mapping by seed-and-vote. *Nucleic Acids Res.* **41**, e108 (2013).
45. Dobin, A. *et al.* STAR: ultrafast universal RNA-seq aligner. *Bioinformatics* **29**, 15–21 (2013).
46. Thorvaldsdóttir, H., Robinson, J.T. & Mesirov, J.P. Integrative Genomics Viewer (IGV): high-performance genomics data visualization and exploration. *Brief. Bioinform.* **14**, 178–192 (2013).
47. Fiume, M., Williams, V., Brook, A. & Brudno, M. Savant: genome browser for high-throughput sequencing data. *Bioinformatics* **26**, 1938–1944 (2010).
48. Fiume, M. *et al.* Savant genome browser 2: visualization and analysis for population-scale genomics. *Nucleic Acids Res.* **40**, 1–7 (2012).
49. Morgan, M. *et al.* ShortRead: a Bioconductor package for input, quality assessment and exploration of high-throughput sequence data. *Bioinformatics* **25**, 2607–2608 (2009).
50. Li, H. *et al.* The sequence alignment/map format and SAMtools. *Bioinformatics* **25**, 2078–2079 (2009).
51. Brooks, A.N. *et al.* Conservation of an RNA regulatory map between *Drosophila* and mammals. *Genome Res.* **21**, 193–202 (2011).
52. Edgar, R., Domrachev, M. & Lash, A.E. Gene expression omnibus: NCBI gene expression and hybridization array data repository. *Nucleic Acids Res.* **30**, 207–210 (2002).
53. Cox, D.R. & Reid, N. Parameter orthogonality and approximate conditional inference. *J. Roy. Stat. Soc. Ser. B Method.* **49**, 1–39 (1987).
54. Dudoit, S., Yang, Y.H., Callow, M.J. & Speed, T.P. Statistical methods for identifying differentially expressed genes in replicated cDNA microarray experiments. *Stat. Sinica* **12**, 111–139 (2002).
55. Bourgon, R., Gentleman, R. & Huber, W. Independent filtering increases detection power for high-throughput experiments. *Proc. Natl. Acad. Sci. USA* **107**, 9546–9551 (2010).
56. Robinson, M.D. & Oshlack, A. A scaling normalization method for differential expression analysis of RNA-seq data. *Genome Biol.* **11**, R25 (2010).
57. Cappiello, C., Francalanci, C. & Pernici, B. Data quality assessment from the user's perspective. *Architecture* **22**, 68–73 (2004).
58. Benjamini, Y. & Hochberg, Y. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *J. Roy. Stat. Soc. Ser. B Methodol.* **57**, 289–300 (1995).
59. Wu, H., Wang, C. & Wu, Z. A new shrinkage estimator for dispersion improves differential expression detection in RNA-seq data. *Biostatistics* **14**, 232–243 (2012).
60. Smyth, G.K. Limma: linear models for microarray data. In *Bioinformatics and Computational Biology Solutions Using R and Bioconductor* (eds. Gentleman, R. *et al.*) 397–420 (Springer, 2005).
61. Nookaew, I. *et al.* A comprehensive comparison of RNA-seq-based transcriptome analysis from reads to differential gene expression and cross-comparison with microarrays: a case study in *Saccharomyces cerevisiae*. *Nucleic Acids Res.* **40**, 10084–10097 (2012).
62. Rapaport, F. *et al.* Comprehensive evaluation of differential expression analysis methods for RNA-seq data <http://arXiv.org/abs/1301.5277v2> (23 January 2013).
63. Hansen, K.D., Irizarry, R.A. & Wu, Z. Removing technical variability in RNA-seq data using conditional quantile normalization. *Biostatistics* **13**, 204–216 (2012).
64. Risso, D., Schwartz, K., Sherlock, G. & Dudoit, S. GC-content normalization for RNA-seq data. *BMC Bioinform.* **12**, 480 (2011).
65. Delhomme, N., Padiou, I., Furlong, E.E. & Steinmetz, L. easyRNASeq: a Bioconductor package for processing RNA-seq data. *Bioinformatics* **28**, 2532–2533 (2012).
66. Leisch, F. Sweave: dynamic generation of statistical reports using literate data analysis. In *Compstat 2002 Proceedings in Computational Statistics* Vol. 69 (eds. Härdle, W. & Rönz, B.) 575–580. Institut für Statistik und Wahrscheinlichkeitstheorie, Technische Universität Wien (Physica Verlag, 2002).